

Microsoft Movie Market Analysis

J.D.Denney

Part 1, Introduction

Microsoft, the client, has tasked our firm with the analysis of film industry data trends and provide insights into their movie making ventures. All analysis and recommendations will be key to their initial efforts at penetrating the market. Let's first pull in the data and take a look at it.

```
In [2]: #import Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
# import sqlite3

# current issues - 1. there are duplicate titles, as many as 4 duplicates for
```

```
In [3]: #initial pull of data with pandas read of title basics csv to get a feel for
dfbscs = pd.read_csv(r"C:\Users\josep\Desktop\CourseWork\phase_1\Phase1\Movie
pd.options.display.float_format = '{:.3f}'.format # to remove the scientific
```

We've imported our first data and formatted any float objects moving forward. What does a quick look at the first five rows of `imdb.title.basics.csv` gives us access to the titles of the movies as well as runtime and start year.

```
In [4]: dfbscs.head()
```

	tconst	primary_title	original_title	start_year	runtime_min
0	tt0063540	Sunghursh	Sunghursh	2013	175.000
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.000
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.000
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh	2018	nan
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.000

```
In [5]: # Let's pull in the other two recommended databases to take a look at what is  
dftrngs = pd.read_csv(r"C:\Users\josep\Desktop\CourseWork\phase_1\Phase1\Mov:  
dfgrss = pd.read_csv(r"C:\Users\josep\Desktop\CourseWork\phase_1\Phase1\Movie
```

In [6]:

```

# And get some info on each of them while taking a look at their size and data
# tconst column will be a key between basics and ratings, while title will be
# of a key between basics and movie gross dataframes
dfbscs.info()
dfrtns.info()
dfgrss.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146144 entries, 0 to 146143
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tconst                 146144 non-null object
1   primary_title          146144 non-null object
2   original_title         146123 non-null object
3   start_year             146144 non-null int64
4   runtime_minutes        114405 non-null float64
5   genres                 140736 non-null object
dtypes: float64(1), int64(1), object(4)
memory usage: 6.7+ MB

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 73856 entries, 0 to 73855
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tconst                 73856 non-null object
1   averagerating          73856 non-null float64
2   numvotes               73856 non-null int64
dtypes: float64(1), int64(1), object(1)
memory usage: 1.7+ MB

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   title                 3387 non-null  object
1   studio                3382 non-null  object
2   domestic_gross        3359 non-null  float64
3   foreign_gross         2037 non-null  object
4   year                  3387 non-null  int64
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB

```

Some of these dataFrames have a lot of rows of data, some. There's a lot going on here, but let's first make sure that we have the right data that will be used to provide insights and make recommendations for Microsoft.

In [7]:

dfbscs.head()

	tconst	primary_title	original_title	start_year	runtime_min
0	tt0063540	Sunghursh	Sunghursh	2013	175.000
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.000
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.000
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh	2018	nan
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.000

In [8]:

dfrtngs.head()

	tconst	averagerating	numvotes
0	tt10356526	8.300	31
1	tt10384606	8.900	559
2	tt1042974	6.400	20
3	tt1043726	4.200	50352
4	tt1060240	6.500	21

In [9]:

dfgrss.head()

	title	studio	domestic_gross	foreign_gross	year
0	Toy Story 3	BV	415000000.000	652000000	2010
1	Alice in Wonderland (2010)	BV	334200000.000	691300000	2010
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.000	664300000	2010
3	Inception	WB	292600000.000	535700000	2010
4	Shrek Forever After	P/DW	238700000.000	513900000	2010

```
In [10]: df_ratings = pd.merge(dfbscs,dfrtngs, on='tconst') # titles without ratings c
df_ratings
# quick check to see if numbers fall in reasonable parameters
print(df_ratings['averagerating'].max()) # = 10
print(df_ratings['averagerating'].min()) # = 1.0
print(df_ratings['start_year'].max()) # = 2019
print(df_ratings['start_year'].min()) # = 2010
# merging data allows for a single consolidated list and removes any titles i
# average rating or numvotes
# can then merge the new df with dfgrss to get gross amounts
# but have to change the column headings on dfbscs first to title from prima

10.0
1.0
2019
2010
```

In [11]:

```

# create a new df, df_clean that can merge on column 'title' after renaming
df_clean = df_ratings.rename(columns = {'primary_title':'title'})
df_clean = df_clean.sort_values(by=['numvotes'],axis=0,ascending=False,ignore
df_clean = df_clean.drop_duplicates(subset='title') # after deleting duplicat
print(len(df_clean))

# check for duplicate titles, verify that "The Door" has multiple entries

# *****want to remove duplicate titles with the least numvotes this is temp
#
# duplicates = df_clean[df_clean.duplicated(subset='title')]
# print(len(duplicates))
# print(duplicates.head())
# for x in df_clean['title']:
#     if x=='The Door':
#         print(x)
# *****

df_ratings_and_gross = pd.merge(df_clean,dfgrss, on='title')
# we now have two sets of data - one 3027 (df_ratings_and_gross) rows long t
# and domestic and foreign gross. The other data set has 73856 records, (df_
# can still provide data on overall ratings.
# There is still some data cleaning to do before we begin visualizing the da
# additionally, one set contains foreign gross which is set to a sting dataty
# to be converted to a float object

# since there are duplicates in the data, I want to sort the df and have the

df_ratings_and_gross = df_ratings_and_gross.sort_values(by=['numvotes'],axis=
# now we can utilize .drop_duplicates on the titles column in order to remov
print(df_ratings_and_gross)
df_ratings_and_gross = df_ratings_and_gross.drop_duplicates(subset='title')
print(df_ratings_and_gross) # note that rows have decreased to 2598 from 302

```

```

69993
      tconst      title      original_title  start_year \
0  tt1375666      Inception      Inception      2010
1  tt1345836  The Dark Knight Rises  The Dark Knight Rises  2012
2  tt0816692      Interstellar      Interstellar      2014
3  tt1853728  Django Unchained  Django Unchained      2012
4  tt0993846  The Wolf of Wall Street  The Wolf of Wall Street  2013
...      ...      ...      ...      ...

```

2594	tt8851190	Red	Red	2018
2595	tt1666555	Anchor Baby	Anchor Baby	2010
2596	tt1692325	Eyes Wide Open	Eyes Wide Open	2010
2597	tt3436064	The Last Station	La última estación	2012
2598	tt2713406	Meerkats	Meerkats	2011

	runtime_minutes	genres	averagerating	numvotes	\
0	148.000	Action,Adventure,Sci-Fi	8.800	1841066	
1	164.000	Action,Thriller	8.400	1387769	
2	169.000	Adventure,Drama,Sci-Fi	8.600	1299334	
3	165.000	Drama,Western	8.400	1211405	
4	180.000	Biography,Crime,Drama	8.200	1035358	
...	
2594	90.000	Drama	8.100	26	
2595	95.000	Drama,Thriller	7.000	25	
2596	110.000	Documentary,History	8.700	17	
2597	90.000	Documentary	7.600	10	
2598	40.000	Documentary	7.400	7	

	studio	domestic_gross	foreign_gross	year
0	WB	292600000.000	535700000	2010
1	WB	448100000.000	636800000	2012
2	Par.	188000000.000	489400000	2014
3	Wein.	162800000.000	262600000	2012
4	Par.	116900000.000	275100000	2013
...
2594	Sum.	90400000.000	108600000	2010
2595	AGF	15800.000	161000	2011
2596	NAV	26300.000	250000	2010
2597	SPC	6600000.000	6900000	2010
2598	NGE	778000.000	482000	2012

[2599 rows x 12 columns]

	tconst	title	original_title	start_year	\
0	tt1375666	Inception	Inception	2010	
1	tt1345836	The Dark Knight Rises	The Dark Knight Rises	2012	
2	tt0816692	Interstellar	Interstellar	2014	
3	tt1853728	Django Unchained	Django Unchained	2012	
4	tt0993846	The Wolf of Wall Street	The Wolf of Wall Street	2013	
...	
2594	tt8851190	Red	Red	2018	
2595	tt1666555	Anchor Baby	Anchor Baby	2010	
2596	tt1692325	Eyes Wide Open	Eyes Wide Open	2010	
2597	tt3436064	The Last Station	La última estación	2012	
2598	tt2713406	Meerkats	Meerkats	2011	

	runtime_minutes	genres	averagerating	numvotes	\
0	148.000	Action,Adventure,Sci-Fi	8.800	1841066	
1	164.000	Action,Thriller	8.400	1387769	
2	169.000	Adventure,Drama,Sci-Fi	8.600	1299334	
3	165.000	Drama,Western	8.400	1211405	
4	180.000	Biography,Crime,Drama	8.200	1035358	
...	
2594	90.000	Drama	8.100	26	
2595	95.000	Drama,Thriller	7.000	25	
2596	110.000	Documentary,History	8.700	17	
2597	90.000	Documentary	7.600	10	

2598		40.000	Documentary	7.400	7
------	--	--------	-------------	-------	---

	studio	domestic_gross	foreign_gross	year
0	WB	292600000.000	535700000	2010
1	WB	448100000.000	636800000	2012
2	Par.	188000000.000	489400000	2014
3	Wein.	162800000.000	262600000	2012
4	Par.	116900000.000	275100000	2013
...
2594	Sum.	90400000.000	108600000	2010
2595	AGF	15800.000	161000	2011
2596	NAV	26300.000	250000	2010
2597	SPC	6600000.000	6900000	2010
2598	NGE	778000.000	482000	2012

[2598 rows x 12 columns]

In [12]:

```

# change foreign gross to a float obj
df_ratings_and_gross['foreign_gross'] = pd.to_numeric(df_ratings_and_gross['foreign_gross'], errors='coerce')
# in this ~3000 records dataset, we need to look for outliers or invalid data
# for max, mins in float data. nothing stands out
print(df_ratings_and_gross['averagerating'].max()) # = 9.2
print(df_ratings_and_gross['averagerating'].min()) # = 1.6
print(df_ratings_and_gross['year'].max()) # = 2018
print(df_ratings_and_gross['year'].min()) # = 2010
print(df_ratings_and_gross['start_year'].max()) # = 2019
print(df_ratings_and_gross['start_year'].min()) # = 2010

```

8.8
1.6
2018
2010
2019
2010

I want to also make sure that all NaN values are ignored when proceeding with the data. Keep in mind that there are NaN values still present in the data, let's move forward and answer some questions about the data and what it means.

Part 2, Question 1

Does a longer runtime indicate a lower average rating? Should Microsoft have a range for their projects?


```
In [13]: # For this initial question, we need to ignore NaN values in average rating.  
# NaN values, there are some movies with high or low ratings but very few votes  
# that the average ratings are higher or lower than they would otherwise be  
# For that reason, I will exclude data with 10 votes or fewer when comparing
```

In [14]:

```

# create new df to call for lmplo
data_df = df_clean[['runtime_minutes', 'averagerating', 'numvotes']]
# filter for numvotes less than 10
df_over_ten = data_df[data_df.numvotes > 10]
# filter for movie outliers with ridiculous lengths
df_over_ten = df_over_ten[df_over_ten.runtime_minutes <= 400]
df_over_ten = df_over_ten.dropna()
# we have removed all rows with a NaN present. Additionally, we can see that
# excluded data that has less than 6 votes that are contributing to the ratio
print(df_over_ten)
# count values to take another look at the data
df_over_ten['runtime_minutes'].value_counts().sort_values()
print(df_over_ten.corr())
# also standard deviation for runtime minutes to figure out a good range for
df_over_ten.mean()
df_over_ten.std()
# 68% of movies are produced within 22 minutes of a 95 minute average.

```


	runtime_minutes	averagerating	numvotes
2387	148.000	8.800	1841066
2241	164.000	8.400	1387769
280	169.000	8.600	1299334
12072	165.000	8.400	1211405
325	143.000	8.100	1183655
...
30670	85.000	6.500	11
30560	45.000	5.900	11
11947	102.000	6.300	11
11502	89.000	7.600	11
34460	240.000	8.500	11

[52456 rows x 3 columns]

	runtime_minutes	averagerating	numvotes
runtime_minutes	1.000	0.018	0.125
averagerating	0.018	1.000	0.066
numvotes	0.125	0.066	1.000

runtime_minutes	22.422
averagerating	1.431
numvotes	35847.867
dtype:	float64

```
In [15]: # drop numvotes column from df_over_five and create an lmplo in seaborn
df_plot1 = df_over_ten.drop(columns=['numvotes'])
# plot lmplo
sns.set(font_scale=2)
sns.set_style("darkgrid")
palette = sns.set_palette("dark")
ax = sns.lmplot(x='runtime_minutes',y='averagerating',data=df_plot1,height=10)
ax.set(xlabel="Runtime in Minutes",ylabel="Average Movie Rating")
plt.title('Average Rating vs. Total Runtime')
plt.show()
```



Question 1 Insights

There is no readily apparent correlation between total runtime and the average rating based on the information in the data tables. Microsoft should focus on making the standard deviation (+/- 22 minutes) of the mean of 95 minutes.

In [16]:

```
df_plot1.corr()
```

as we can see, there is no correlation between runtime_minutes and averagerating

	runtime_minutes	averagerating
runtime_minutes	1.000	0.018
averagerating	0.018	1.000

Part 3, Question 2

What can we learn by looking at the correlation between total gross and the rating of a movie? Put slightly differently, does a higher quality movie demand more money at the BO?

In [17]:

just as a reminder, here is what the table we will be working with for question 2


```
df_ratings_and_gross.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2598 entries, 0 to 2598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tconst                 2598 non-null   object
1   title                  2598 non-null   object
2   original_title         2598 non-null   object
3   start_year             2598 non-null   int64
4   runtime_minutes        2598 non-null   float64
5   genres                 2598 non-null   object
6   averagerating          2598 non-null   float64
7   numvotes               2598 non-null   int64
8   studio                 2595 non-null   object
9   domestic_gross         2581 non-null   float64
10  foreign_gross          1554 non-null   float64
11  year                   2598 non-null   int64
dtypes: float64(4), int64(3), object(5)
memory usage: 263.9+ KB
```

In [18]:

```
# For this question, we will look at a graphical representation of the average
# plotted against the total gross (domestic_gross + foreign_gross)
# sum foreign and domestic into new column, ignoring NaN values of which there are
df_ratings_and_gross['total_gross'] = df_ratings_and_gross['domestic_gross'] +
df_ratings_and_gross['foreign_gross'] # but how many of the total gross are 0 after adding a new
# domestic/foreign columns
df_ratings_and_gross.isna().sum()
df_ratings_and_gross['total_gross'].value_counts()
df_ratings_and_gross['total_gross'] = df_ratings_and_gross['total_gross'] / 1000000
# great, we can see that total gross has 0 elements that show NaN, as predicted
# If a movie had NaN in both domestic and foreign gross then it is showing as NaN
# after checking for value counts == 0 we can see that we have only good data
# assuming that all NaN's were meant to be 0's. For domestic gross, that is less
# be unusual but not impossible for a movie to make money abroad but not in its
```

```
In [19]: # plot graph of average rating vs total gross
sns.set(font_scale=2)
sns.set_style("darkgrid")
palette = sns.set_palette("dark")
ax = sns.lmplot(x='averagerating',y=('total_gross'),data=df_ratings_and_gross)
ax.set(xlabel="Average Rating",ylabel="Total Gross in Billions")
plt.ticklabel_format(style='plain',axis='y')
plt.title('Total Gross Scatter Plot')
plt.show()
# issue here is y axis label with 1e9 top left, make it prettier
df_plot2 = df_ratings_and_gross.drop(columns=['start_year','runtime_minutes',
```



Question 2 Insights

As seen in the graph above and the correlation table below, there is only a weak positive correlation between average rating and the total gross of a movie. The higher quality a movie is,

gross more at the box office. We still haven't hit on a truly valuable insight for

In [20]: `df_plot2.corr()`

	averagerating	domestic_gross	foreign_gross	total_gross
averagerating	1.000	0.154	0.186	0.147
domestic_gross	0.154	1.000	0.826	0.920
foreign_gross	0.186	0.826	1.000	0.978
total_gross	0.147	0.920	0.978	1.000

Part 4, Question 3

What genre of films gross the most at the box office (BO)? What other insights can you get from this?

In [21]: `# We will need to break out the genres and list them separately for all films`
`# are multiple genres listed.`
`# df_ratings_and_gross has total gross in it from our previous exercise`
`len(df_ratings_and_gross['genres'].unique())`
`# Let's separate these out`

296

In [22]: `df_ratings_and_gross.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2598 entries, 0 to 2598
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tconst                 2598 non-null  object
1   title                  2598 non-null  object
2   original_title         2598 non-null  object
3   start_year             2598 non-null  int64
4   runtime_minutes        2598 non-null  float64
5   genres                 2598 non-null  object
6   averagerating          2598 non-null  float64
7   numvotes               2598 non-null  int64
8   studio                 2595 non-null  object
9   domestic_gross         2581 non-null  float64
10  foreign_gross          1554 non-null  float64
11  year                   2598 non-null  int64
12  total_gross            2598 non-null  float64
dtypes: float64(5), int64(3), object(5)
memory usage: 284.2+ KB
```

```
In [23]: # need to separate genres out into a list and count total occurrences
# see above for list of unique genre combinations - there are too many to put
# histogram with.
#genres_list = df_ratings_and_gross['genres'].dropna().str.lower().str.split(' ')
#g_list = []

#for genre_list in genres_list:
#    try:
#        g_list.extend(genre_list)
#    except:
#        continue
```

```
In [24]: # Raf dummify genres example
# df_example = pd.get_dummies(df_ratings_and_gross['genres'].str.split(' ').c
```

Lets group our data by genres, including those where genre is null. We can then look at what the averages of the data can show us.


In [25]:

```
df_genres = df_ratings_and_gross.groupby(df_ratings_and_gross['genres']).mean()
# Let's graph the top 20 of
# this for total_gross and single out some genres (and genre combinations) that
# are successful.
df_genres2 = df_genres.sort_values('total_gross',axis=0,ascending=False).reset_index()
df_genres2.head(15) # we can see here the top 15 grossing genres or combinations
```

	index	genres	start_year	runtime_minutes	average_rating	numvotes
0	106	Adventure,Fantasy	2013.333	139.667	7.167	375770.333
1	101	Adventure,Drama,Sci-Fi	2014.500	156.500	8.300	989725.000
2	11	Action,Adventure,Sci-Fi	2013.978	131.370	6.833	428425.111
3	30	Action,Comedy,Mystery	2018.000	121.000	6.100	1250.000
4	148	Biography,Drama,Musical	2017.000	105.000	7.600	199663.000
5	8	Action,Adventure,Fantasy	2015.065	117.903	6.287	250437.652
6	110	Adventure,Mystery,Sci-Fi	2012.000	124.000	7.000	538720.000
7	12	Action,Adventure,Thriller	2013.750	125.000	6.550	224555.000
8	104	Adventure,Family,Fantasy	2015.125	120.000	6.113	164708.625
9	72	Action,Sci-Fi	2014.000	113.000	7.900	546284.000
10	75	Adventure,Animation,Comedy	2014.200	94.453	6.439	107169.000
11	96	Adventure,Drama,Fantasy	2013.400	118.800	6.840	254748.000
12	68	Action,Mystery,Sci-Fi	2014.000	113.000	6.800	387038.000
13	45	Action,Drama,Family	2010.500	133.500	6.650	214967.500
14	4	Action,Adventure,Comedy	2015.355	112.806	6.342	187106.800

In [26]:

```
# graph represents the average per genre rating and the average per genre total gross
sns.set(font_scale=2)
sns.set_style("darkgrid")
palette = sns.set_palette("dark")
ax = sns.lmplot(x='averagerating',y='total_gross',data=df_genres2,height=10,palette=palette)
ax.set(xlabel="Average Rating",ylabel="Total Gross in $ Billions")
plt.ticklabel_format(style='plain',axis='y')
plt.title('Average Gross Scatter Plot by Genre')
plt.show()
df_genres2.corr()
```




	index	start_year	runtime_minutes	averagerating	numvotes	domestic_gross
index	1.000	0.011	-0.165	0.069	-0.155	-0.236
start_year	0.011	1.000	0.098	0.169	-0.021	0.048
runtime_minutes	-0.165	0.098	1.000	0.052	0.282	0.220
averagerating	0.069	0.169	0.052	1.000	0.189	0.050
numvotes	-0.155	-0.021	0.282	0.189	1.000	0.788
domestic_gross	-0.236	0.048	0.220	0.050	0.788	1.000
foreign_gross	-0.311	0.253	0.324	0.108	0.646	0.763
year	0.018	0.954	0.076	0.215	-0.088	-0.026
total_gross	-0.298	0.107	0.266	0.075	0.761	0.906


In [27]:

```

# one scatter plot for the highest 15 grossing genres and have average rating
sns.set(font_scale=2)
sns.set_style("darkgrid")
palette = sns.set_palette("dark")
ax = sns.lmplot(x='averagerating',y='total_gross',data=df_genres2.head(20),hu
ax.set(xlabel="Average Rating",ylabel="Total Gross in $ Billions")
plt.ticklabel_format(style='plain',axis='y')
plt.title('Average Gross Scatter Plot by Genre, Top 20')
plt.show()

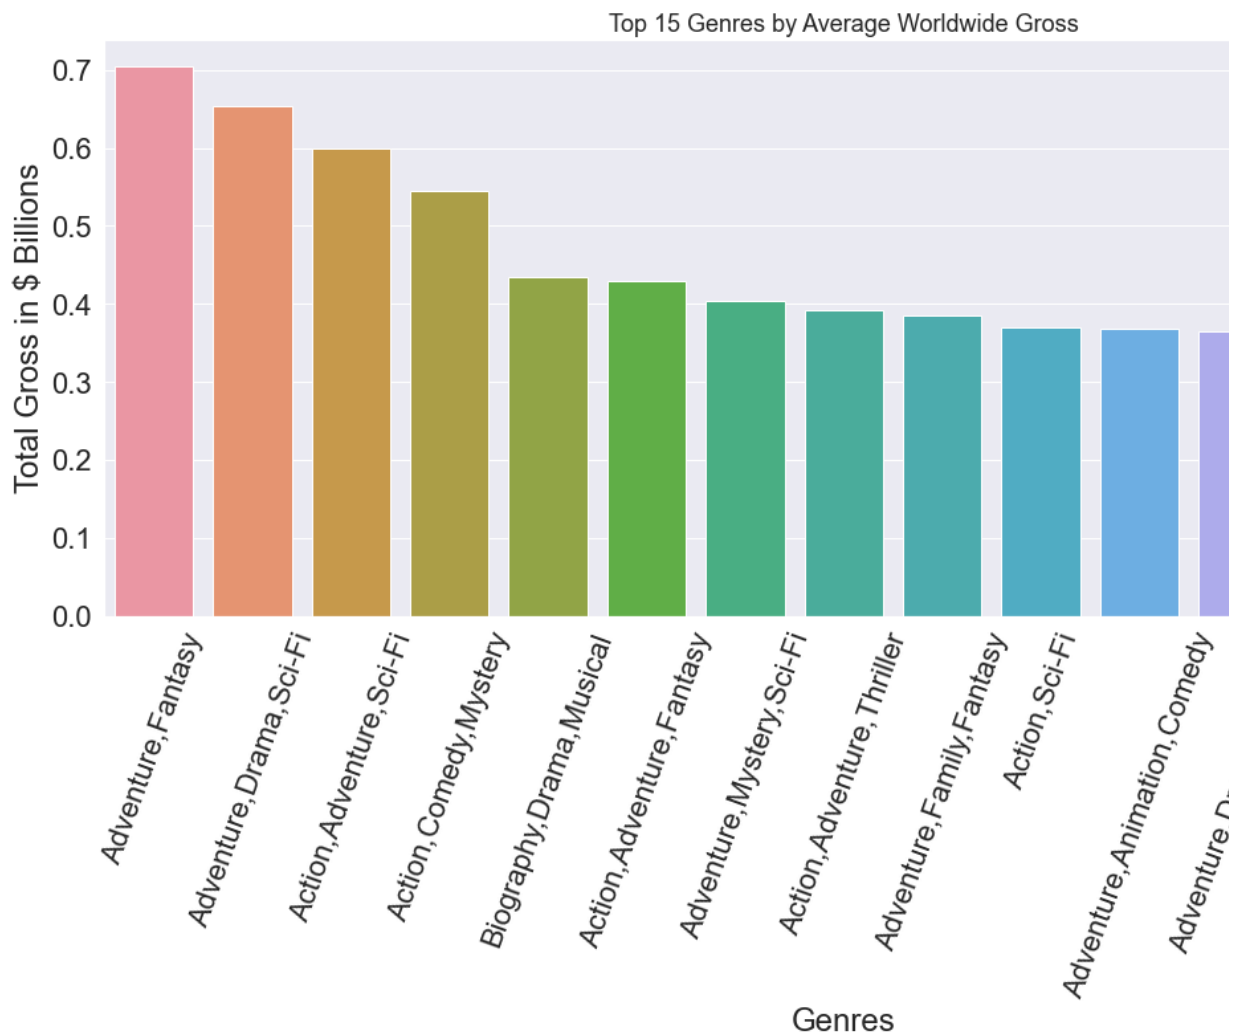
```





In [28]:

```
# one bar chart with genres on x axis, total gross on y axis
fig_dims = (20,8)
fig, ax = plt.subplots(figsize=fig_dims)
sns.set(font_scale=1.5)
sns.set_style("darkgrid")
palette = sns.set_palette("dark")
ax = sns.barplot(x=df_genres2.head(15)['genres'], y=df_genres2.head(15)['total_gross'])
ax.set(xlabel="Genres",ylabel="Total Gross in $ Billions")
plt.ticklabel_format(style='plain',axis='y')
plt.xticks(rotation=70)
plt.title('Top 15 Genres by Average Worldwide Gross')
plt.show()
```



Question 3 Insights

Adventure, Action, Fantasy, and Sci-Fi movies are all heavily featured as the most popular genres. Microsoft should begin by focusing on some of these more popular genres.

market, even if there is low ROI up front. Creating good cornerstone content foundation for creating profits in the future.

Part 5, Question 4

What Genres of film provides the highest ROI?

```
In [29]: dfmovie = pd.read_csv(r"C:\Users\josep\Desktop\CourseWork\phase_1\Phase1\Micro
dfbudgets = pd.read_csv(r"C:\Users\josep\Desktop\CourseWork\phase_1\Phase1\Micro
```

```
In [30]: dfbudgets = dfbudgets.rename(columns = {'movie':'title'})
dfbudgets_merge = pd.merge(df_ratings_and_gross,dfbudgets, on='title') # merge
# # movies with extremely low production budgets are going to create some outliers
dfbudgets_2 = dfbudgets_merge.drop(columns=['domestic_gross_y', 'domestic_gross_x'])
dfbudgets_2['production_budget'] = dfbudgets_2['production_budget'].str.replace('M', '')
dfbudgets_2['production_budget'] = dfbudgets_2['production_budget'].str.replace('B', '')
dfbudgets_2['production_budget'] = dfbudgets_2['production_budget'].astype('float')
avebudget = dfbudgets_2['production_budget'].mean()
stdevbudget = dfbudgets_2['production_budget'].std()
# # I want to remove the outlier data outside of 2 standard deviations with pandas
# # miraculous movies like Paranormal Activity that had a huge ROI. This should be removed
lower_bound = avebudget - (stdevbudget * 2)
higher_bound = avebudget + (stdevbudget * 2)
```

In [31]:

```
dfbudgets_2 = dfbudgets_2[dfbudgets_2.production_budget >= lower_bound]
dfbudgets_2 = dfbudgets_2[dfbudgets_2.production_budget <= higher_bound]
print(dfbudgets_2.info()) # check to see if 5% of data removed - verified
dfbudgets_2['production_budget'] = dfbudgets_2['production_budget']/100000000
dfbudgets_2['worldwide_gross'] = dfbudgets_2['worldwide_gross'].str.replace('
dfbudgets_2['worldwide_gross'] = dfbudgets_2['worldwide_gross'].str.replace('
dfbudgets_2['worldwide_gross'] = dfbudgets_2['worldwide_gross'].astype('float
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1087 entries, 0 to 1165
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tconst                 1087 non-null  object
1   title                 1087 non-null  object
2   original_title        1087 non-null  object
3   start_year            1087 non-null  int64
4   runtime_minutes       1087 non-null  float64
5   genres                1087 non-null  object
6   averagerating         1087 non-null  float64
7   numvotes              1087 non-null  int64
8   studio                1087 non-null  object
9   year                  1087 non-null  int64
10  total_gross            1087 non-null  float64
11  id                     1087 non-null  int64
12  release_date           1087 non-null  object
13  production_budget      1087 non-null  float64
14  worldwide_gross        1087 non-null  object
dtypes: float64(4), int64(4), object(7)
memory usage: 135.9+ KB
None
```

In [32]:

```
# for the purposes of the question we will be looking at total_gross divided

dfbudgets_2 = dfbudgets_2.groupby(dfbudgets_2['genres']).mean().reset_index('
dfbudgets_2['ROI'] = (dfbudgets_2['total_gross']/(dfbudgets_2['production_bu
dfbudgets_2_sorted = dfbudgets_2.sort_values('ROI',axis=0,ascending=False)).re
```

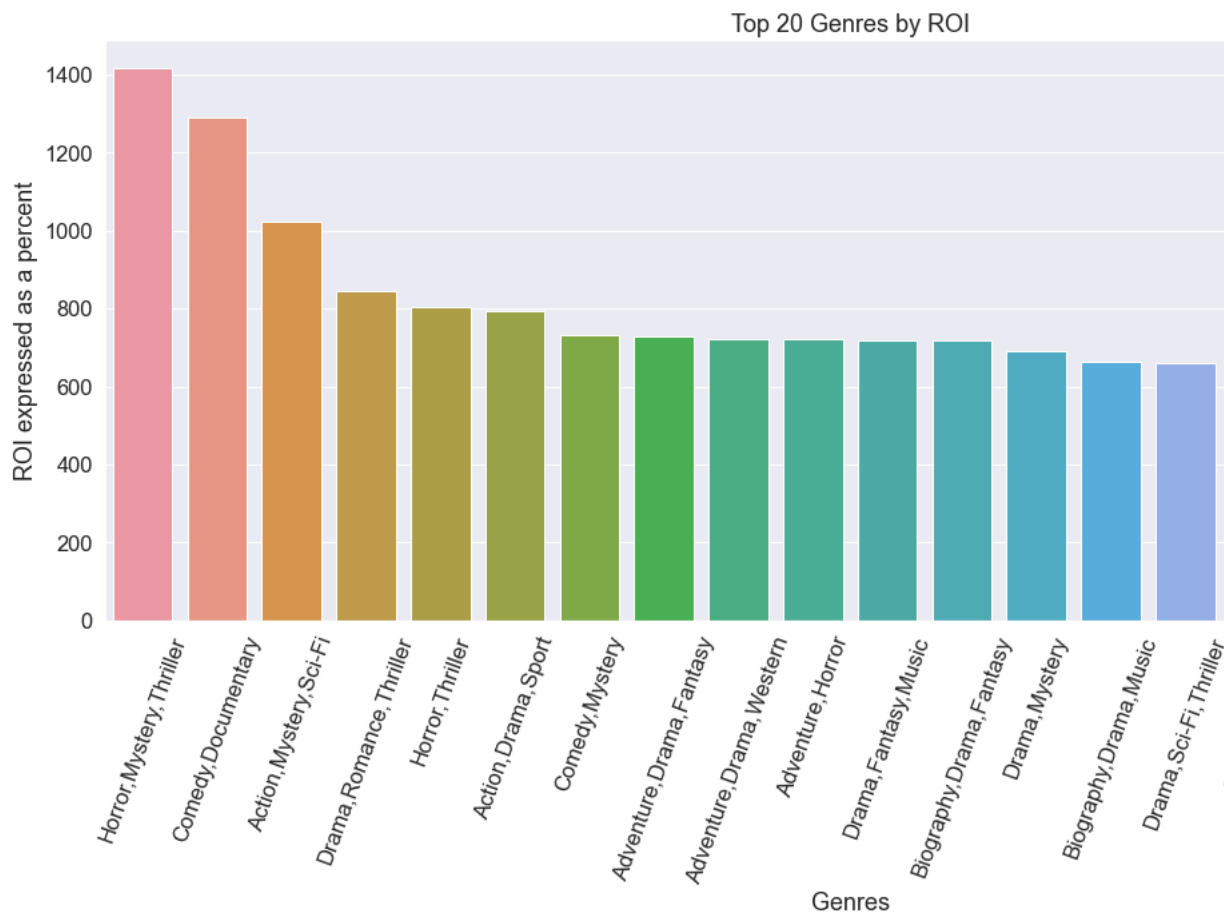

In [33]:

dfbudgets_2_sorted.head(20)

	index	genres	start_year	runtime_minutes	averagerating	numvotes
0	180	Horror,Mystery,Thriller	2014.192	93.769	5.500	93775.654
1	99	Comedy,Documentary	2013.000	75.000	6.800	7880.000
2	49	Action,Mystery,Sci-Fi	2014.000	113.000	6.800	387038.000
3	165	Drama,Romance,Thriller	2015.333	115.000	5.100	140024.667
4	182	Horror,Thriller	2014.176	98.706	5.571	90904.059
5	39	Action,Drama,Sport	2016.000	170.000	7.100	33371.000
6	120	Comedy,Mystery	2011.000	102.000	6.500	432800.000
7	69	Adventure,Drama,Fantasy	2012.750	116.750	6.875	283280.250
8	73	Adventure,Drama,Western	2010.000	110.000	7.600	284034.000
9	76	Adventure,Horror	2013.000	100.000	5.300	33239.000
10	145	Drama,Fantasy,Music	2014.000	107.000	6.800	107625.000
11	86	Biography,Drama,Fantasy	2018.000	99.000	5.400	24351.000
12	159	Drama,Mystery	2010.000	100.000	7.200	127751.000
13	88	Biography,Drama,Music	2014.600	123.400	6.680	87972.400
14	168	Drama,Sci-Fi,Thriller	2011.500	92.500	7.050	397175.000
15	113	Comedy,Fantasy	2012.500	97.500	5.750	188033.500
16	153	Drama,Horror,Mystery	2014.900	101.100	6.130	83364.000
17	33	Action,Drama,Family	2010.333	135.667	6.500	192112.000
18	47	Action,Horror,Sci-Fi	2015.333	100.500	5.733	80642.500
19	179	Horror,Mystery,Sci-Fi	2013.000	84.500	5.200	44551.000

In [35]:

```
# one bar chart with genres on x axis, total gross on y axis
fig_dims = (20,8)
fig, ax = plt.subplots(figsize=fig_dims)
sns.set(font_scale=1.5)
sns.set_style("darkgrid")
palette = sns.set_palette("dark")
ax = sns.barplot(x=dfbudgets_2_sorted.head(20)['genres'], y=dfbudgets_2_sorted['ROI'])
ax.set(xlabel="Genres",ylabel="ROI expressed as a percent")
plt.ticklabel_format(style='plain',axis='y')
plt.xticks(rotation=70)
plt.title('Top 20 Genres by ROI')
plt.show()
```



Question 4 Insights

There are some extremely high ROI's for certain combinations of genres. At foundation of quality content and a user base supportive of that content, we utilize less than the average for production budgets while diversifying content finding material to develop in the Horror, Thriller, Mystery, and Sci-Fi will prove in the long run. The worst thing Microsoft could do is throw money at projects that underperform without first establishing quality original content.

In []: